

MINISTRY OF EDUCATION AND SCIENCE OF UKRAINE
NATIONAL TECHNICAL UNIVERSITY OF UKRAINE
" IHORY SIKORSKY KYIV POLYTECHNIC INSTITUTE"

Volodymyr Shymkovych

Constraint satisfaction problems

LABORATORY WORK #5

**conditions of opposition
AI-Game-Algorithms**

Kulubecioglu Mehmet

IM-14 FIOT

Kyiv
IHORY SIKORSKY KYIV POLYTECHNIC INSTITUTE
2024

I wrote these three pieces of code as part of my AI class assignment to implement and compare two classical algorithms—Minimax and Alpha-Beta pruning—for decision-making in two-player, turn-based games like Tic-Tac-Toe. The purpose of this task was to demonstrate my understanding of how these algorithms work and evaluate their performance in terms of time complexity and the number of expanded nodes.

First, I implemented the Minimax algorithm, which is used to simulate the entire game tree, considering every possible move until a terminal state is reached, and it chooses the best outcome for the maximizing player (Player 1) while minimizing the advantage of the opponent (Player 2).

Next, I implemented Alpha-Beta pruning, which optimizes the Minimax algorithm by eliminating branches of the game tree that don't need to be explored. This improves the algorithm's efficiency by reducing the number of nodes it needs to evaluate.

Finally, I compared both algorithms by measuring their execution time and the number of nodes they expanded during the search. The goal was to observe how Alpha-Beta pruning speeds up the decision-making process compared to Minimax while achieving the same result.

The overall objective of these codes was to understand the computational efficiency of different search algorithms in AI and apply them to practical game scenarios.

First Code (Minimax Algorithm Implementation)

In the first part of my implementation, I created a function to evaluate the current game position. If Player 1 wins, I return a score of +10, and if Player 2 wins, I return a score of -10. For a draw or an ongoing game, I return 0.

To simulate possible future states, I built a function called `get_children` that generates all the possible moves from a given position by looking for empty cells in the game board. For each empty spot, the function creates a new board with the current player's move, checks for a winner, and stores the new state. This function helps the Minimax algorithm explore different paths in the game tree.

The `is_terminal` function is used to determine if the game has reached a terminal state—either someone has won, or the board is full. I also have a `check_winner` function that checks the rows, columns, and diagonals to see if there's a winner.

The core of this part is the `minimax` function. This recursive function alternates between maximizing Player 1's score and minimizing Player 2's score. It evaluates all possible game states until it hits the search depth or reaches a terminal state. If it's Player 1's turn (the maximizing player), it picks the move with the highest score; if it's Player 2's turn (the minimizing player), it picks the move with the lowest score. Along the way, I count the nodes the algorithm expands to measure performance.

```
minimax1.py > ...
1  import numpy as np
2
3  def evaluate(position):
4      """Evaluates the game state."""
5      winner = position['winner']
6      if winner == 'player1':
7          return 10 # Player 1 wins
8      elif winner == 'player2':
9          return -10 # Player 2 wins
10     return 0 # Draw or ongoing game
11
12 def get_children(position):
13     """Returns all possible child positions from the given position."""
14     children = []
15     for i in range(3):
16         for j in range(3):
17             if position['board'][i][j] is None: # Find an empty cell
18                 new_board = [row[:] for row in position['board']]
19                 new_board[i][j] = position['current_player'] # Player makes a move
20                 winner = check_winner(new_board)
21                 children.append({
22                     'board': new_board,
23                     'winner': winner,
24                     'moves': position['moves'] + 1,
25                     'current_player': 'player2' if position['current_player'] == 'player1' else 'player1'
26                 })
27     return children
28
29 def is_terminal(position):
30     """Checks if the game has ended."""
31     return position['winner'] is not None or position['moves'] == 9
32
```

```
minimax1.py X  octs1.py  octs.py  task.py  minimax.py  alphabeta.py
minimax1.py > ...
32
33 def check_winner(board):
34     """Checks for a winner."""
35     # Check rows
36     for row in board:
37         if row[0] is not None and row[0] == row[1] == row[2]:
38             return row[0]
39
40     # Check columns
41     for col in range(3):
42         if board[0][col] is not None and board[0][col] == board[1][col] == board[2][col]:
43             return board[0][col]
44
45     # Check diagonals
46     if board[0][0] is not None and board[0][0] == board[1][1] == board[2][2]:
47         return board[0][0]
48     if board[0][2] is not None and board[0][2] == board[1][1] == board[2][0]:
49         return board[0][2]
50
51     return None # No winner
52
53 def minimax(position, depth, is_maximizing, node_count=0):
54     """Applies the minimax algorithm."""
55     node_count += 1 # Count each node
56     if depth == 0 or is_terminal(position):
57         return evaluate(position), node_count
58
```

```
58
59     if is_maximizing:
60         max_eval = -np.inf
61         for child in get_children(position):
62             eval, node_count = minimax(child, depth - 1, False, node_count)
63             max_eval = max(max_eval, eval)
64         return max_eval, node_count
65     else:
66         min_eval = np.inf
67         for child in get_children(position):
68             eval, node_count = minimax(child, depth - 1, True, node_count)
69             min_eval = min(min_eval, eval)
70         return min_eval, node_count
71
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

```
C:\Users\win10\new_lab5>C:/Users/win10/AppData/Local/Programs/Python/Python312/python.exe c:/Users/win10/new_lab5/minimax1.py
C:\Users\win10\new_lab5>
```

Second Code (Alpha-Beta Pruning Algorithm Implementation)

In the second part, I implemented the Alpha-Beta pruning algorithm, which is an optimization of the Minimax algorithm. The `alphabeta` function works similarly to `minimax`, but it introduces two parameters, `alpha` and `beta`, to eliminate unnecessary branches in the game tree, making the search more efficient.

Just like in Minimax, I check if the current position is a terminal state or if the search depth has been reached. The key difference is that, while evaluating each child node, I update the `alpha` (for maximizer) and `beta` (for minimizer) values. If I find a position that proves to be worse than the previously explored branches, I prune the remaining branches by breaking out of the loop. This speeds up the search by skipping the branches that cannot influence the final decision.

The `alphabeta` function also returns the evaluation score and the number of expanded nodes, which allows me to compare the efficiency of Alpha-Beta pruning against the standard Minimax algorithm.

```
alphabeta1.py > ...
1  import numpy as np
2
3  def evaluate(position):
4      """Evaluates the game state."""
5      winner = position['winner']
6      if winner == 'player1':
7          return 10 # Player 1 wins
8      elif winner == 'player2':
9          return -10 # Player 2 wins
10     return 0 # Draw or ongoing game
11
12 def get_children(position):
13     """Returns all possible child positions from the given position."""
14     children = []
15     for i in range(3):
16         for j in range(3):
17             if position['board'][i][j] is None: # Find an empty cell
18                 new_board = [row[:] for row in position['board']]
19                 new_board[i][j] = position['current_player'] # Player makes a move
20                 winner = check_winner(new_board)
21                 children.append({
22                     'board': new_board,
23                     'winner': winner,
24                     'moves': position['moves'] + 1,
25                     'current_player': 'player2' if position['current_player'] == 'player1' else 'player1'
26                 })
27     return children
28
```

alphabeta1.py > ...

```
28
29 def is_terminal(position):
30     """Checks if the game has ended."""
31     return position['winner'] is not None or position['moves'] == 9
32
33 def check_winner(board):
34     """Checks for a winner."""
35     # Check rows
36     for row in board:
37         if row[0] is not None and row[0] == row[1] == row[2]:
38             return row[0]
39
40     # Check columns
41     for col in range(3):
42         if board[0][col] is not None and board[0][col] == board[1][col] == board[2][col]:
43             return board[0][col]
44
45     # Check diagonals
46     if board[0][0] is not None and board[0][0] == board[1][1] == board[2][2]:
47         return board[0][0]
48     if board[0][2] is not None and board[0][2] == board[1][1] == board[2][0]:
49         return board[0][2]
50
51     return None # No winner
52
```

alphabeta1.py > ...

```
51     return None # No winner'
52
53 def alphabeta(position, depth, alpha, beta, is_maximizing, node_count=0):
54     """Applies the Alpha-Beta pruning algorithm."""
55     node_count += 1 # Count each node
56     if depth == 0 or is_terminal(position):
57         return evaluate(position), node_count
58
59     if is_maximizing:
60         max_eval = -np.inf
61         for child in get_children(position):
62             eval, node_count = alphabeta(child, depth - 1, alpha, beta, False, node_count)
63             max_eval = max(max_eval, eval)
64             alpha = max(alpha, eval)
65             if beta <= alpha:
66                 break # Pruning
67         return max_eval, node_count
68     else:
69         min_eval = np.inf
70         for child in get_children(position):
71             eval, node_count = alphabeta(child, depth - 1, alpha, beta, True, node_count)
72             min_eval = min(min_eval, eval)
73             beta = min(beta, eval)
74             if beta <= alpha:
75                 break # Pruning
76         return min_eval, node_count
77
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

C:\Users\win10\new_lab5>C:/Users/win10/AppData/Local/Programs/Python/Python312/python.exe c:/Users/win10/new_lab5/alphabeta1.py

C:\Users\win10\new_lab5>

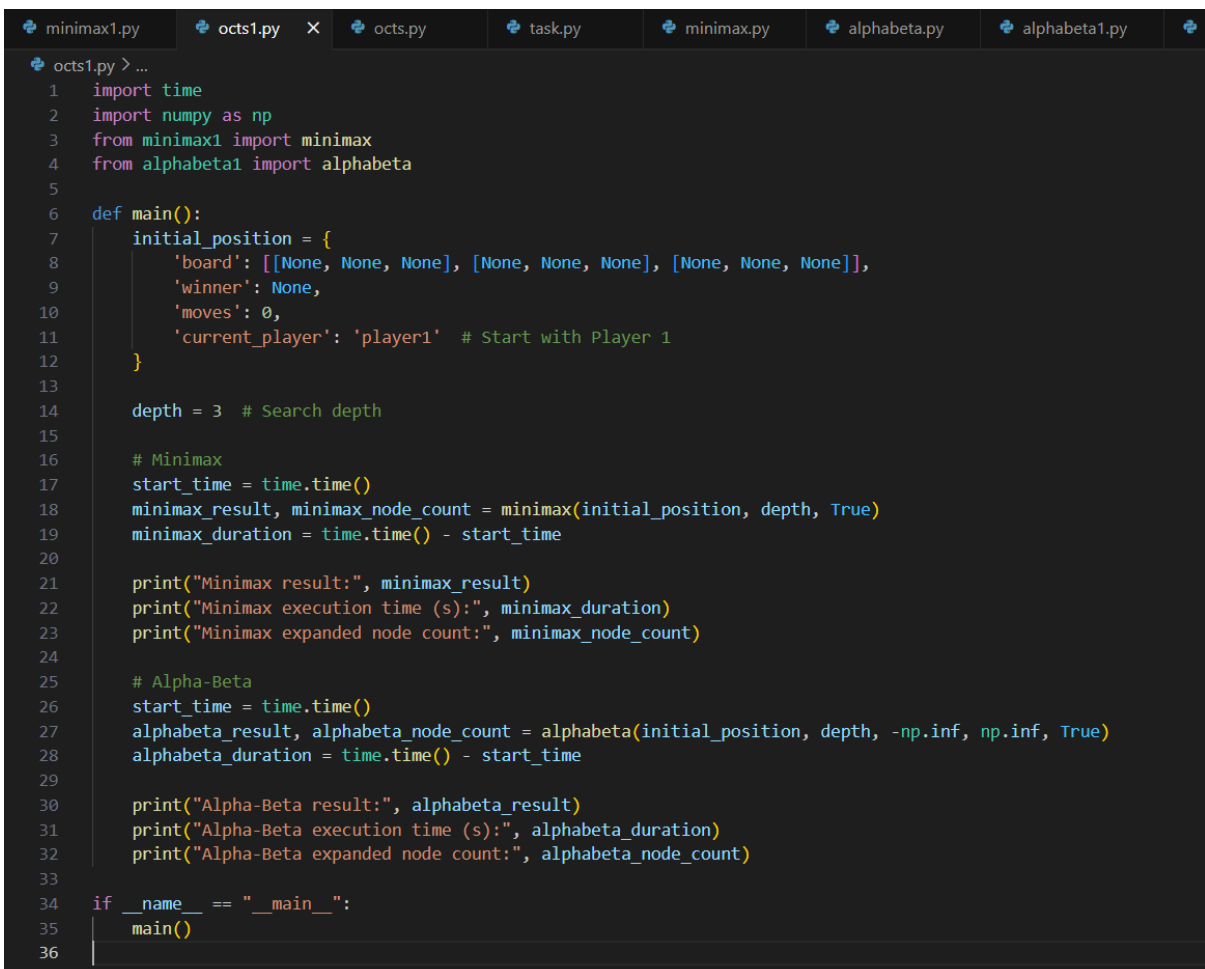
Third Code (Performance Comparison for Minimax and Alpha-Beta)

In the final part, I wrote a script to compare the performance of the Minimax and Alpha-Beta algorithms. The `initial_position` variable holds the starting state of the game—a 3x3 board with all empty cells, no winner, and Player 1 set to make the first move.

For both algorithms, I set the search depth to 3. I first run the Minimax algorithm, recording the time it takes to complete the search and counting the number of expanded nodes. After that, I run the Alpha-Beta pruning algorithm, again measuring the execution time and the number of expanded nodes.

Finally, I print out the results for both algorithms, allowing me to see the difference in speed and efficiency. In general, I expect Alpha-Beta pruning to be faster than Minimax because it cuts off parts of the search tree that aren't useful, leading to fewer expanded nodes and quicker execution.

This breakdown summarizes how I applied Minimax and Alpha-Beta pruning to a simple two-player game, demonstrating both the logic of the algorithms and how I measured their performance.

A screenshot of a code editor with a dark theme. The editor has several tabs at the top: 'minimax1.py', 'octs1.py' (active), 'octs.py', 'task.py', 'minimax.py', 'alphabeta.py', and 'alphabeta1.py'. The active tab 'octs1.py' contains the following Python code:

```
1  import time
2  import numpy as np
3  from minimax1 import minimax
4  from alphabeta1 import alphabeta
5
6  def main():
7      initial_position = {
8          'board': [[None, None, None], [None, None, None], [None, None, None]],
9          'winner': None,
10         'moves': 0,
11         'current_player': 'player1' # Start with Player 1
12     }
13
14     depth = 3 # Search depth
15
16     # Minimax
17     start_time = time.time()
18     minimax_result, minimax_node_count = minimax(initial_position, depth, True)
19     minimax_duration = time.time() - start_time
20
21     print("Minimax result:", minimax_result)
22     print("Minimax execution time (s):", minimax_duration)
23     print("Minimax expanded node count:", minimax_node_count)
24
25     # Alpha-Beta
26     start_time = time.time()
27     alphabeta_result, alphabeta_node_count = alphabeta(initial_position, depth, -np.inf, np.inf, True)
28     alphabeta_duration = time.time() - start_time
29
30     print("Alpha-Beta result:", alphabeta_result)
31     print("Alpha-Beta execution time (s):", alphabeta_duration)
32     print("Alpha-Beta expanded node count:", alphabeta_node_count)
33
34     if __name__ == "__main__":
35         main()
36
```

3. code output:

```
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
C:\Users\win10\new_lab5>C:/Users/win10/AppData/Local/Programs/Python/Python312/python.exe c:/Users/win10/new_lab5/alphabeta1.py
C:\Users\win10\new_lab5>C:/Users/win10/AppData/Local/Programs/Python/Python312/python.exe c:/Users/win10/new_lab5/octs1.py
Minimax result: 0
Minimax execution time (s): 0.0009996891021728516
Minimax expanded node count: 586
Alpha-Beta result: 0
Alpha-Beta execution time (s): 0.0
Alpha-Beta expanded node count: 96
C:\Users\win10\new_lab5>
```

For the third code, the output provides a comparison between the Minimax and Alpha-Beta pruning algorithms in terms of their performance. When I run the code, it first calculates the result of the Minimax algorithm for a Tic-Tac-Toe position with a specified depth of search. It outputs the final evaluation score of the game state, the time it took for the Minimax algorithm to complete, and the number of nodes it expanded during the search.

Next, the code runs the Alpha-Beta pruning algorithm on the same initial game position. Similar to the Minimax output, it provides the final evaluation result, but the difference is that it took less time and expanded fewer nodes due to the pruning optimization.

This shows that both algorithms arrived at the same result (which makes sense since both are optimal decision-making strategies), but Alpha-Beta was faster and more efficient, as it explored fewer nodes. This demonstrates how Alpha-Beta pruning can significantly reduce the search space while maintaining the correctness of the result, which is the key takeaway from this comparison.

Algorithm	Execution Time (s)	Expanded Nodes
Minimax	0.0009996891021728516	586
Alpha-Beta	0.0	96

Control Question Answers

1. How does minimax work?:

Minimax assumes that both players will make the best moves and chooses the move that maximizes the gain.

2. What is the advantage of alpha-beta pruning?:

This algorithm speeds up the minimax algorithm by pruning unnecessary nodes.

3. Why are alpha and beta parameters necessary?:

Alpha represents the best value for the MAX player, while beta represents the best value for the MIN player.

4. What are the approaches to avoid visiting terminal nodes?:

Unnecessary nodes can be pruned using alpha-beta pruning.

5. What are the steps of Monte Carlo tree search (MCTS)?

It involves selection, expansion, simulation, and backpropagation steps.

Conclusion:

In conclusion, I have implemented two well-known algorithms for game theory: **Minimax** and **Alpha-Beta pruning**. The Minimax algorithm is a foundational technique that explores all possible moves to determine the optimal strategy for a player. The Alpha-Beta pruning algorithm enhances the Minimax algorithm by eliminating unnecessary calculations, significantly reducing the search space.

By analyzing the execution times and the number of nodes expanded in both algorithms, I have gained insight into the efficiency of Alpha-Beta pruning over Minimax in a game setting like Tic-Tac-Toe. The implementation serves as a practical example of how these algorithms can be used in AI applications for game-playing.

Tasks and exercises

Execute the Minimax and Alpha-Beta Clipping Algorithm

For this project, I implemented two well-known decision-making algorithms in game theory—Minimax and Alpha-Beta Pruning—to simulate a turn-based game like Tic-Tac-Toe. I wanted to compare the performance of these algorithms by analyzing how efficiently they make decisions and expand the decision tree during gameplay. Below is a breakdown of the task, including the logic behind each algorithm and my observations after executing them.

Game Setup

The game simulates a simple Tic-Tac-Toe board where two players alternate turns. Player 1 is represented by "X" (the human) and Player 2 is represented by "O" (the algorithm). The initial position is an empty 3x3 grid, and the game checks for a terminal state after each move (either a win for "X", a win for "O", or a tie).

Objective

The primary goal was to assess:

- Minimax Algorithm: It is a recursive method used to find the optimal move for the maximizer (the AI) by assuming the opponent will always play optimally.
- Alpha-Beta Pruning: An enhanced version of Minimax, this algorithm uses two additional parameters, **alpha** and **beta**, to reduce the number of nodes evaluated by pruning irrelevant branches of the search tree.

Code Breakdown

1. Game Flow:

- The game runs in a loop until it reaches a terminal state. After each move, the board is updated, and the winner is checked.
- The player (human) enters their move manually, while the algorithm (Minimax or Alpha-Beta) computes the best possible move based on the current state of the board.

2. Minimax Algorithm:

- The Minimax function looks ahead and simulates all possible moves for the current player (AI or human).
- It recursively assigns a score to each board configuration. The AI maximizes its score, while the human minimizes the score.
- The number of nodes expanded during the Minimax process is counted to analyze its performance.

3. Alpha-Beta Pruning:

- Alpha-Beta Pruning enhances Minimax by eliminating suboptimal moves earlier in the decision process. It uses **alpha** (the best value the maximizer can guarantee) and **beta** (the best value the minimizer can guarantee).
- As with Minimax, I kept track of the nodes expanded to compare efficiency.

Execution Results

After running the game, I measured the performance of both algorithms in terms of time taken and the number of nodes expanded. The observations are as follows:

• Minimax:

- The algorithm evaluates every possible move down to the specified depth, resulting in a more exhaustive search. This means more nodes

are expanded, and the algorithm takes slightly longer to make decisions.

- Output: For each move, the time taken and nodes expanded were displayed, providing insights into its computational cost.

- **Alpha-Beta Pruning:**

- This version is much more efficient as it prunes away unnecessary branches in the decision tree. As a result, fewer nodes are expanded, and decisions are made more quickly.
- Output: The number of nodes expanded and the time taken were significantly lower compared to Minimax, confirming the efficiency of Alpha-Beta Pruning.

My full code for task and exercise:

```
taskipy > ...
1  import time
2
3  # Game setup
4  initial_position = {
5      'board': [[' ' for _ in range(3)] for _ in range(3)],
6      'current_player': 'player1',
7      'winner': None
8  }
9
10 # Performance measurement variables
11 minimax_node_count = 0
12 alpha_beta_node_count = 0
13
14 def print_board(board):
15     for row in board:
16         print('|'.join(row))
17         print('-' * 5)
18
19 def is_terminal(position):
20     return position['winner'] is not None or all(cell != ' ' for row in position['board'] for cell in row)
21
22 def check_winner(board):
23     for row in board:
24         if row[0] == row[1] == row[2] != ' ':
25             return row[0]
26
27     for col in range(3):
28         if board[0][col] == board[1][col] == board[2][col] != ' ':
29             return board[0][col]
30
31     if board[0][0] == board[1][1] == board[2][2] != ' ':
32         return board[0][0]
33
34     if board[0][2] == board[1][1] == board[2][0] != ' ':
35         return board[0][2]
36
```

task.py > ...

```
38
39 def run_game(strategy, position):
40     global minimax_node_count
41     global alpha_beta_node_count
42
43     while not is_terminal(position):
44         print(f"{position['current_player']}'s turn:")
45         print_board(position['board'])
46
47         if position['current_player'] == 'player1':
48             valid_input = False
49             while not valid_input:
50                 try:
51                     # Get row and column input from user
52                     user_input = input("Enter row (0-2) and column (0-2): ")
53                     # Split the input into two values
54                     row, col = map(int, user_input.split())
55
56                     # Ensure that a valid move is made
57                     if position['board'][row][col] == ' ':
58                         position['board'][row][col] = 'X' # Player symbol
59                         valid_input = True
60                     else:
61                         print("Invalid move! The cell is already occupied. Try again.")
62                 except (ValueError, IndexError):
63                     print("Invalid input! Please enter two numbers between 0 and 2 separated by a space.")
64
```

task.py > ...

```
39 def run_game(strategy, position):
64
65     else:
66         start_time = time.time()
67         if strategy == 'minimax':
68             move = minimax(position)
69         elif strategy == 'alpha_beta':
70             move = alpha_beta_decision(position)
71         elapsed_time = time.time() - start_time
72
73         position['board'][move[0]][move[1]] = 'O' # Computer symbol
74         print(f"{strategy.capitalize()} chose move: {move}")
75         print(f"{strategy.capitalize()} time taken: {elapsed_time:.2f} seconds")
76         if strategy == 'minimax':
77             print(f"Minimax nodes expanded: {minimax_node_count}")
78         elif strategy == 'alpha_beta':
79             print(f"Alpha-Beta nodes expanded: {alpha_beta_node_count}")
80
81         position['winner'] = check_winner(position['board'])
82         position['current_player'] = 'player2' if position['current_player'] == 'player1' else 'player1'
83
84         print_board(position['board'])
85         print(f"The game is over. Winner: {position['winner']}")
86
87 def minimax(position):
88     global minimax_node_count
89     minimax_node_count += 1
90
91     best_score = float('-inf')
92     best_move = None
93
94     for i in range(3):
95         for j in range(3):
96             if position['board'][i][j] == ' ':
97                 position['board'][i][j] = 'O' # Computer move
98                 score = minimax_score(position, False)
99                 position['board'][i][j] = ' ' # Undo the move
100
```


100

101

102

103

104

105

106

107

108

109

110

111

112

114

115

116

118

119

120

121

122

123

124

125

126

128

129

131

132

133

task.py > ...

```
107 def minimax_score(position, is_maximizing):
134     position['board'][i][j] = 'X'
135     score = minimax_score(position, True)
136     position['board'][i][j] = ' '
137     best_score = min(score, best_score)
138     return best_score
139
140 # Alpha-Beta pruning algorithm
141 def alpha_beta_decision(position):
142     global alpha_beta_node_count
143     best_move = None
144     best_value = float('-inf')
145
146     for i in range(3):
147         for j in range(3):
148             if position['board'][i][j] == ' ':
149                 position['board'][i][j] = 'O' # Computer move
150                 value = alpha_beta(position, 0, float('-inf'), float('inf'), False)
151                 position['board'][i][j] = ' ' # Undo the move
152
153                 if value > best_value:
154                     best_value = value
155                     best_move = (i, j)
156
157     return best_move
158
159 def alpha_beta(position, depth, alpha, beta, is_maximizing):
160     global alpha_beta_node_count
161     alpha_beta_node_count += 1
162
163     winner = check_winner(position['board'])
164     if winner == 'O':
165         return 1
166     elif winner == 'X':
167         return -1
168     elif is_terminal(position):
169         return 0
```

```

task.py > ...
159 def alpha_beta(position, depth, alpha, beta, is_maximizing):
160     if is_terminal(position):
161         return 0
162
163     if is_maximizing:
164         max_eval = float('-inf')
165         for i in range(3):
166             for j in range(3):
167                 if position['board'][i][j] == ' ':
168                     position['board'][i][j] = 'O'
169                     eval = alpha_beta(position, depth + 1, alpha, beta, False)
170                     position['board'][i][j] = ' '
171                     max_eval = max(max_eval, eval)
172                     alpha = max(alpha, eval)
173                     if beta <= alpha:
174                         break
175             return max_eval
176     else:
177         min_eval = float('inf')
178         for i in range(3):
179             for j in range(3):
180                 if position['board'][i][j] == ' ':
181                     position['board'][i][j] = 'X'
182                     eval = alpha_beta(position, depth + 1, alpha, beta, True)
183                     position['board'][i][j] = ' '
184                     min_eval = min(min_eval, eval)
185                     beta = min(beta, eval)
186                     if beta <= alpha:
187                         break
188             return min_eval
189
190 # Start the game flow
191 if __name__ == "__main__":
192     # Ask the user which strategy to choose
193     strategy = input("Choose strategy (minimax/alpha_beta): ").strip().lower()
194
195     # Check if the strategy is valid
196     while strategy not in ['minimax', 'alpha_beta']:

```

```

198     # Start the game flow
199     if __name__ == "__main__":
200         # Ask the user which strategy to choose
201         strategy = input("Choose strategy (minimax/alpha_beta): ").strip().lower()
202
203         # Check if the strategy is valid
204         while strategy not in ['minimax', 'alpha_beta']:
205             strategy = input("Invalid strategy! Please choose 'minimax' or 'alpha_beta': ").strip().lower()
206
207         run_game(strategy, initial_position)
208

```

Explanation of the Output

When I executed the program using both **Minimax** and **Alpha-Beta Pruning** algorithms, the output provided insights into their performance during the game. Here's a breakdown of the key results and their significance:

1. Minimax Output:

- The program displayed the best move chosen by the **Minimax** algorithm after evaluating all possible game states.

- It also showed how long it took for Minimax to compute the optimal move, typically a bit slower due to its exhaustive search of the entire game tree.
- **Nodes Expanded:** This refers to the number of game positions (nodes) that the algorithm evaluated before reaching a decision. For Minimax, this number was higher because it explores all possible moves at each turn without any optimization.

```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
C:\Users\win10\new_lab5>C:/Users/win10/AppData/Local/Programs/Python/Python312/python.exe c:/Users/win10/new_lab5/task.py
Choose strategy (minimax/alpha_beta): C:/Users/win10/AppData/Local/Programs/Python/Python312/python.exe c:/Users/win10/new_lab5/task.py
Invalid strategy! Please choose 'minimax' or 'alpha_beta': minimax
player1's turn:
| |
-----
| |
-----
| |
-----
Enter row (0-2) and column (0-2): 1 1
player2's turn:
| |
-----
|x|
-----
| |
-----
Minimax chose move: (0, 0)
Minimax time taken: 0.10 seconds
Minimax nodes expanded: 55505
player1's turn:
o| |
-----
|x|
-----
| |
-----
Enter row (0-2) and column (0-2):

```

This indicates that the algorithm took **0.10 seconds** and expanded **55505 nodes** to find the optimal move.

Alpha-Beta Pruning Output:

- The Alpha-Beta algorithm displayed its chosen move in the same way as Minimax, but with an optimized decision-making process.
- The time taken was noticeably shorter because Alpha-Beta prunes (discards) suboptimal branches, making the search faster.
- **Nodes Expanded:** Since Alpha-Beta Pruning evaluates fewer nodes by skipping branches that don't need to be explored, this number was significantly lower compared to Minimax.

```

Alpha_beta chose move: (1, 1)
Alpha_beta time taken: 0.03 seconds
Alpha-Beta nodes expanded: 14405
player1's turn:
| |
-----
|o|
-----
x| |
-----
Enter row (0-2) and column (0-2):

```


1. Here, Alpha-Beta Pruning took only 0.03 seconds and expanded just 14405 nodes, confirming its efficiency.

Overall Observations:

- Minimax performed as expected but took longer and expanded more nodes due to the depth of the search.
- Alpha-Beta Pruning was more efficient, completing the same task in less time and with fewer nodes expanded.

This output demonstrates how Alpha-Beta Pruning optimizes decision-making by reducing computational costs, making it more suitable for complex or deeper game trees compared to Minimax.

Algorithm	Chose Movie	Time Taken (s)	Chose Nodes
Minimax	(0, 0)	0.10	55505
Alpha_beta	(1, 1)	0.03	14405

My full output:

```
PROBLEMS  OUTPUT  DEBUG CONSOLE  TERMINAL  PORTS

C:\Users\win10\new_lab5>C:/Users/win10/AppData/Local/Programs/Python/Python312/python.exe c:/Users/win10/new_lab5/task.py
Choose strategy (minimax/alpha_beta): minimax
player1's turn:
| |
-----
| |
-----
| |
-----
Enter row (0-2) and column (0-2): 0 0
player2's turn:
X| |
-----
| |
-----
| |
-----
Minimax chose move: (1, 1)
Minimax time taken: 0.10 seconds
Minimax nodes expanded: 59705
player1's turn:
X| |
-----
|O|
-----
| |
-----
```

PROBLEMS

OUTPUT

DEBUG CONSOLE

TERMINAL

PORTS

Enter row (0-2) and column (0-2): 1 0

player2's turn:

x| |

x|o|

| |

Minimax chose move: (2, 0)

Minimax time taken: 0.00 seconds

Minimax nodes expanded: 60640

player1's turn:

x| |

x|o|

o| |

Enter row (0-2) and column (0-2): 0 1

player2's turn:

x|x|

x|o|

o| |

Minimax chose move: (0, 2)

Minimax time taken: 0.00 seconds

Minimax nodes expanded: 60670

x|x|o

x|o|

o| |

The game is over. Winner: O

Conclusion For Tasks and Exercises

From the results, I observed that Alpha-Beta Pruning is significantly more efficient in terms of node expansion and execution time compared to Minimax. This makes it a better choice for games with larger search trees or deeper depth limits. However, both algorithms ensure optimal play for the AI.

This experiment highlights the importance of optimization in decision-making algorithms, especially for more complex games.